



Singularity Project

Yet another operating system?

by Miroslav Rudišin

March 2006



About project

■ Who?

- Microsoft Research OS group
 - Galen Hunt, James Larus, ...

■ What?

- building more dependable software platform

■ Why?

- insufficient robustness, reliability, security

Dependability

The notion of dependability, defined as the ***trustworthiness*** of a computing system which allows reliance to be justifiably placed on the ***service*** it delivers, enables these various concerns to be subsumed within a ***single*** conceptual ***framework***.

Dependability thus includes as special cases such attributes as reliability, availability, safety, security.



The principle

Everything is under control
(should be)

Software-isolated processes

- SIP = closed object space
 - depends on language safety
 - memory independence invariant
- Communication between SIPs
 - sending messages over channels
 - transfers ownership of data
- *Yay, 1979 again? Mainframes?*

Extensibility

- no dynamic code loading
- SIP = encapsulation
 - on failure system frees SIPs resources and notifies communication partners
 - simply way to isolate and discard corrupted data
- no dynamic code generation
- **compile-time reflection**
 - similar to macros, aspects, multi-stage programming...



Application abstraction

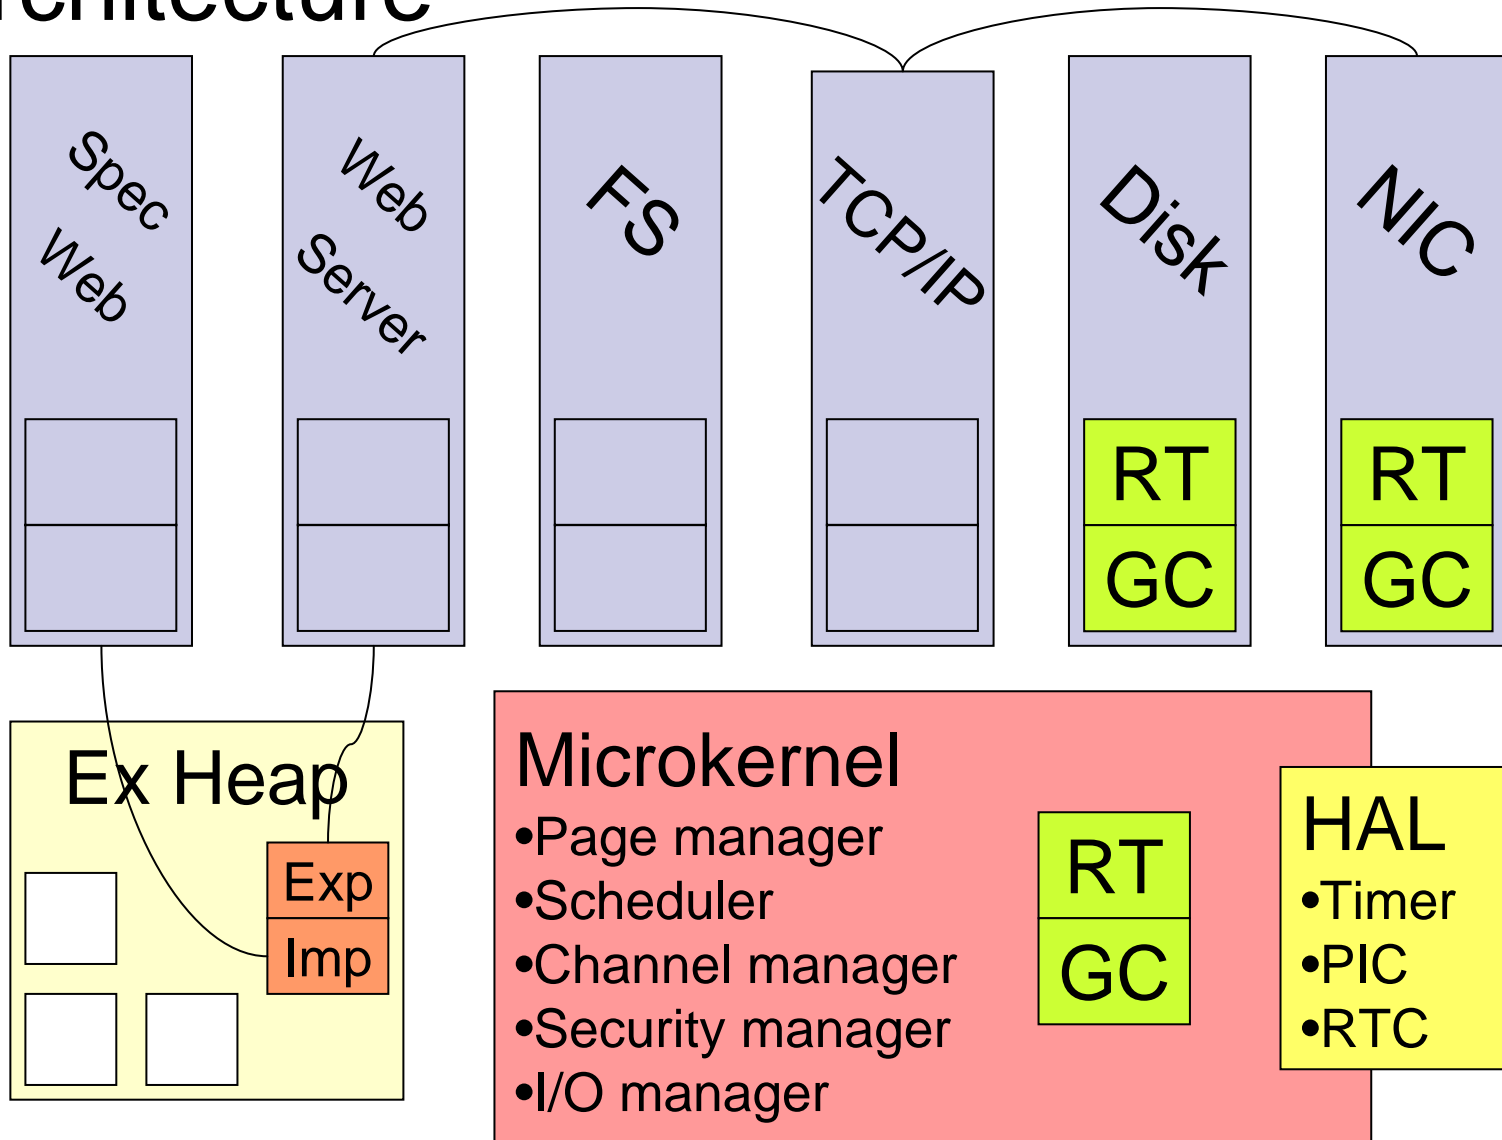
- application = manifest & resources
- manifest describes desired state
- unified Singularity installer
- system is aware of all relationships



Key contributions

- SIP = verified safe code
- consistent extension model
- run-time verification of execution and communication
- language and compiler
- OS = safe language run-time system

Architecture



Code in Singularity

- C# ... Spec# ... **Sing#**
 - verification of type-safety, null pointers, exceptions, contract between functions and also protocol-level interactions between components
- Compiled to safe **MSIL**
- Compiled to x86 code by **Bartok** compiler
 - *interpreted/JIT kernel would be insane (Inferno OS, JX)*
- **Trusted code** (HAL, Kernel, parts of RT system, GC)
- assembler, C++, unsafe C#

Kernel

- privileged system component
- own garbage collected object space
- provides Application Binary Interface (ABI)
 - static kernel methods
 - maintains system-wide **state isolation invariant**
 - ABI versioning – *Microsoft.Singularity.V1.Threads*
- compile-time replaceable scheduler
- exports synchronization constructs – **handle table**
 - used for coordination threads within process

Processes

- single virtual address space
- created by allocating memory sufficient for image
- started by kernel by executing trusted startup code
 - stack & memory pages initialization
- **memory independence invariant**
- stack management (per thread)
 - linked stack (non-contiguous segments)
- *Everything runs in kernel! No context switches. ... False.*

Garbage collection

- prevents memory deallocation errors
- every process is garbage collected (kernel too)
- no universal garbage collector (GC)
- each SIP has its own GC
- 5 selectable GC for each SIP
- independent scheduling of GCs (no cross-object ptrs)
- ABI calls to kernel are executed on thread's stack
 - solved by marking that area - GC can skip marked frames

Run-time system

- each SIP can have completely different RT system
- allows customization
 - sequential code without threads
 - specific allocation strategies such as preallocation
 - ...
- own memory layout, garbage collection algorithms, libraries
- *Can SIP select scheduling policy for its threads? Probably no.*

Exchange Heap

- holds data passed between processes
- reference counts to track usage of **regions**
- process accesses region through structure **allocation**
- strict ownership that maintains memory isolations
- **tracked resources** – via Sing# abstraction **TRef**
 - blocking **Acquire** of ownership, **Release**
- allocating memory on ExHeap, req. **Exchangeable type**
 - R In ExHeap a; /* exchangeable type */*
 - R * In ExHeap pr; /* pointer to R */*
 - R [] In ExHeap pv; /* vector of R */*

Channels

- SIPs communicating by sending messages over channels
- bidirectional, behaviourly typed **contract**
- messages are tagged data in Exchange Heap
- send is asynchronous
- receive synchronously blocks
- **ownership invariant**
 - prevents sharing
 - helps static program analysis
 - provides message-passing semantics

Channel contract

- **Sing#**
- **contract C1 {**
 - in message** Request(int x) **requires** x>0;
 - out message** Reply(int y);
 - out message** Error();
 - state** Start: Request?
 - > (Reply! or Error!)
 - > Start;**}**
- C1.NewChannel(**out** Imp, **out** Exp)

Endpoints

- Pair of endpoints: **Imp**, **Exp**

- C1.Imp {
 void SendRequest(**int** x);
 void RecvReply(**out int** y);
 void RecvError();
}

C1.Exp {
 void RecvRequest(**out int** x);
 void SendReply(**int** y);
 void SendError();
}

Switch-Receive Statement

```
void M (C1.Imp a, C1.Imp b) {  
    switch receive {  
        case a.Reply(x) && b.Reply(y):  
            Console.WriteLine("Both replies {0} and{1}", x, y);  
            break;  
        case a.Error():  
            Console.WriteLine("Error reply on a");  
            break;  
        case b.Error():  
            Console.WriteLine("Error reply on b");  
            break;  
        case a.ChannelClosed():  
            Console.WriteLine("Channel a is closed");  
            break;  
    }  
}
```

Compile-Time Reflection (CTR)

- generators are written in Sing# as transforms
- basic idea ... place-holders are expanded by generator
- usage scenario:
 - device driver describes its resource requirements
 - startup code can be generated from this description
- CTR transform may be part of trusted base, so it can generate trusted code

CTR Example

```
transform DriverTransform
where $IoRangeType: IoRange {
    class $DriverCategory: DriverCategoryDeclaration {
        [$IoRangeAttribute(*)]
        $IoRangeType $$ioranges;
        public readonly static $DriverCategory Values;
        generate static $DriverCategory() {
            Values = new $DriverCategory();
        }
        implement private $DriverCategory() {
            IoConfig config = IoConfig.GetConfig();
            Tracing.Log(Tracing.Debug, "Config: {0}", config.ToPrint());
            forall ($cindex = 0; $f in $$ioranges; $cindex++) {
                $f = ($f.$IoRangeType) config.DynamicRanges[$cindex];
            }
        }
    }
}
```

CTR Example (cont.)

```
internal class Sb16Resources: DriverCategoryDeclaration {  
    [IoPortRange(0, Default = 0x0220,Length = 0x10)]  
    internal readonly IoPortRange basePorts;  
    [IoPortRange(1, Default = 0x0380,Length = 0x10)]  
    internal readonly IoPortRange gamePorts;  
    internal readonly static Sb16Resources Values;  
    reflective private Sb16Resources();  
}
```

```
class SB16Resources {  
    ...  
    static Sb16Resources() { Values = new Sb16Resources(); }  
    private SB16Resources() {  
        IoConfig config = IoConfig.GetConfig();  
        Tracing.Log(Tracing.Debug, "Config: {0}", config.ToPrint());  
        basePorts = (IoPortRange) config.DynamicRanges[0];  
        gamePorts = (IoPortRange) config.DynamicRanges[1];  
    }  
}
```

Verification

- three-stage process
 - Sing# - type safety, ownership rules, protocol
 - Singularity verifier on MSIL code
 - (not yet) Back-end compiler produces typed assembly
 - enables run-time checks by operating system
- mostly with static code analysis
- small impact on performance

Singularity System

- I/O System – 3 layers: HAL, I/O manager, device drivers
- HAL: IoPorts, IoDMA, IoIrq, IoMemory
- device drivers binded by theirs manifests (metadata)
- maintains **3 device driver (DD) invariants**:
 - never installs DD conflicting with other DD
 - never starts DD with conflicting or missing resource
 - DD cannot access unspecified resources

Singularity Namespace

- single, uniform namespace for all services
- UNIX-like mountpoints
- examples:

/hardware/devices

/hardware/drivers

/filesystems/ntfs

/tcp/128.0.0.1/80

/fs/foo/bar

/apps/ms/word

/users/fred

Security

- shape of namespace
- enforcing access by managing channels of application according to its manifest
- discoverable process identity at the peer side of channel
- trace of process execution history
- compound principals:
 - /sys/login @ /users/fred + /apps/ms/word
- access control expressions (ACE)
- restriction on sending some messages on channel
- lending identities

Proposed changes on HW

- memory protection for DMA transfers
 - only unsafe aspect of driver-device interface
- support for segmented stacks
- simpler memory protection for trusted base
 - everything runs in RING0

Performance - microbenchmarks

- AMD Athlon 64 3000+, nForce 4 chipset, 1GB RAM, 7200 SATA, Gigabit NIC
- Cost of basic operations:

	Cost (CPU Cycles)			
	Singularity	FreeBSD	Linux	Windows
Read cycle counter	8	6	6	2
ABI call	87	878	437	627
Thread yield	394	911	906	753
2 thread wait-set ping pong	1,207	4,707	4,041	1,658
2 message ping pong	1,452	13,304	5,797	6,344
Create and start process	300,000	1,032,000	719,000	5,376,000

Disk I/O benchmarks

- random & sequential disk reads and writes
- 1000 operations / 512MB
- Singularity - disk driver SIP, channels – zero copy
- other OSes by system calls
- different blocksizes
- all systems has comparable performance (6% diff)
- Conclusion: Singularity has comparable disk I/O performance

SPECweb99 benchmarks

- Cassini open-source C# web server ported to SIPs
- Singularity: 91 ops/sec
- MS Windows 2003 - IIS: 761 ops/sec
- Singularity – instability under heavy load & FS bottleneck
- Response time: Singularity at 23 ops/sec was 322 ms/op
- Network stack does not appear to be bottleneck

Visions & plans

- better performance
- checking for liveness/deadlocks
- specification of actions when component fails
- distributed Singularity
- integration of heterogenous execution environment
 - .NET, EJB, CCM, DCOM, ...
- ...

```

Disassembly
Offset:
Retrieving information...

Command
System.String: /pci/02/00/1011/0009/20/0a00
System.String: /pnp/PNP0800
System.String: /pnp/PNP0C04
System.String: /pnp/PNP0501
System.String: /pnp/PNP0501
System.String: /pnp/PNP0700
System.String: /pnp/PNP0400
System.String: /pnp/PNP0C02
System.String: /pnp/PNP0C02
Devices associated but not initialized:
Devices initialized:
System.String: Microsoft.Singularity.Driver
System.String: Microsoft.Singularity.Driver
System.String: Microsoft.Singularity.Driver
System.String: Microsoft.Singularity.Driver
System.String: Microsoft.Singularity.Driver
System.String: Microsoft.Singularity.Driver
System.String: Microsoft.Singularity.Driver
Stopping thread 0x70fc39c at index 0x4
Cleaning up after thread 4
Thread.Cleanup calling ISchedulerThread.Clea
Process.Join <end>
Created new thread process.
Process.Start
Process.Join <begin>
[T:0x70fc71c]Created thread 0x70fc71c at ind
#GCEVENT Thread 0x5: StopTheWorld 0x4
Finalizer called for Value=0x4.
Finalizer called for Value=0x3.
Finalizer called for Value=0x2.
Finalizer called for Value=0x1.
#GCEVENT Thread 0x5: ResumeTheWorld 0x4
Stopping thread 0x70fc71c at index 0x5
Cleaning up after thread 5
Thread.Cleanup calling ISchedulerThread.Clea
Process.Join <end>
Unmapped key: 0x84
Unmapped key: 0xfc
Unmapped key: 0x84
Unmapped key: 0xfc

```

Singularity S3Trio64 Driver

Running Sing# Shell.

Type 'help' to get a list of valid commands.

Singularity> pnp

Drivers:

/pci/03/00/5333/8811:

/pnp/PNP0800:

/pnp/PNP0100:

/pnp/PNP0303:

/pnp/PNP0A03:

/pnp/PNP0B00:

Devices found but not associated:

/pnp/080: /pnp/PNP0803/tBA03B0

/pnp/081: /pnp/PNP082F/tBA2FB0

/pnp/100: /pnp/PNP0C01

/pnp/102: /pnp/PNP0200

/pnp/106: /pnp/PNP0F13

/pnp/107/0000/0000/0000: /pci/06/00/8086/7192/03/0000/0000

/pnp/107/0000/0007/0000: /pci/06/01/8086/7110/01/0000/0000

/pnp/107/0000/0007/0001: /pci/01/01/8086/7111/01/0000/0000

/pnp/107/0000/0007/0003: /pci/06/80/8086/7113/02/0000/0000

/pnp/107/0000/000a/0000: /pci/02/00/1011/0009/20/0a00/2114

/pnp/108: /pnp/PNP0800

/pnp/109: /pnp/PNP0C04

/pnp/10a: /pnp/PNP0501

/pnp/10b: /pnp/PNP0501

/pnp/10c: /pnp/PNP0700

/pnp/10d: /pnp/PNP0400

/pnp/10e: /pnp/PNP0C02

/pnp/10f: /pnp/PNP0C02

Devices associated but not initialized:

Devices initialized:

/pnp/0: Microsoft.Singularity.Drivers.PnpBios

/pnp/101: Microsoft.Singularity.Drivers.Pic

/pnp/103: Microsoft.Singularity.Drivers.Timer8254

/pnp/104: Microsoft.Singularity.Drivers.RTClock

/pnp/105: Microsoft.Singularity.Drivers.Keyboard8042

/pnp/107: Microsoft.Singularity.Drivers.Pci.PciBus

/pnp/107/0000/0000/0000: Microsoft.Singularity.Drivers.S3Trio64

Singularity> collect

Linked Data: [4 items]

Value = 4 [Next] 3, 28672 bytes payload

Value = 3 [Next] 2, 27648 bytes payload

Value = 2 [Next] 1, 26624 bytes payload

Value = 1 [Next] 4, 25600 bytes payload

Collecting garbage [before heap: 1413120 bytes]

Collecting garbage [after heap: 1089536 bytes]

Singularity>

The End.

References:

Technical Report (2005-135)

<http://research.microsoft.com/os/singularity/>

Singularity Revisited (video)

<http://channel9.msdn.com/Showpost.aspx?postid=141858>

Jeff Darcy's short review of Singularity

<http://pl.atyp.us/wordpress/?p=991>

Article and discussion on Slashdot

<http://slashdot.org/article.pl?sid=05/11/03/1744230>